

# SlotMarketSQL

Zack Horton, Virginia Maguire, Tanner Street, and Yutong Zhang  
Group A26

## **Problem**

- Models for stock price and volatility forecasts are crucial for day traders.
- However, many traders do not have the skill set to build and interface with these models.

## **Solution**

- Forecast daily closing prices
- Forecast daily volatilities
- Chatbot for easy info access

### Data:

- Historical closing prices (Yahoo Finance)

### Scope:

- 2022-02-22 to 2024-02-21
- Stocks in the S&P 500

SlotMarketSQL helps  
day traders make  
informed decisions

## SlotMarketSQL

Enter your question into the following textbox:

Prompt:

Run query

In the background, it  
references aggregated  
model outputs

Ticker	Stock	current_price	forecasted_price	forecasted_price_change	percent_change	current_volatility	forecasted_volatility	forecasted_volatility_change	model_error_last_week	model_mda_last_week
xyl	xylem inc.	124.099998	111.671234	-12.428764	-10.01512	0.007417	0.015811	113.160105	110.922851	42.857143
yum	yum! brands	133.949997	128.791183	-5.158813	-3.851298	0.004823	0.010888	125.77327	126.736209	57.142857
zbra	zebra technologies	271.929993	261.986816	-9.943176	-3.656521	0.011846	0.028328	139.128409	247.513789	57.142857
zbh	zimmer biomet	124.980003	120.568626	-4.411377	-3.529666	0.000231	0.014593	6228.137813	120.531616	42.857143
zts	zoetis	186.550003	181.352402	-5.197601	-2.786171	0.002227	0.017044	665.448184	182.485736	28.571429

SlotMarketSQL helps  
day traders make  
informed decisions

## SlotMarketSQL

Enter your question into the following textbox:

Prompt:

Run query

In the background, it  
references aggregated  
model outputs

LSTM model  
outputs

Ticker	Stock	current_price	forecasted_price	forecasted_price_change	percent_change	current_volatility	forecasted_volatility	forecasted_volatility_change	model_error_last_week	model_mda_last_week
xyl	xylem inc.	124.099998	111.671234	-12.428764	-10.01512	0.007417	0.015811	113.160105	110.922851	42.857143
yum	yum! brands	133.949997	128.791183	-5.158813	-3.851298	0.004823	0.010888	125.77327	126.736209	57.142857
zbra	zebra technologies	271.929993	261.986816	-9.943176	-3.656521	0.011846	0.028328	139.128409	247.513789	57.142857
zbh	zimmer biomet	124.980003	120.568626	-4.411377	-3.529666	0.000231	0.014593	6228.137813	120.531616	42.857143
zts	zoetis	186.550003	181.352402	-5.197601	-2.786171	0.002227	0.017044	665.448184	182.485736	28.571429

Garch model  
outputs

LSTM performance  
metrics

# LSTM Overview

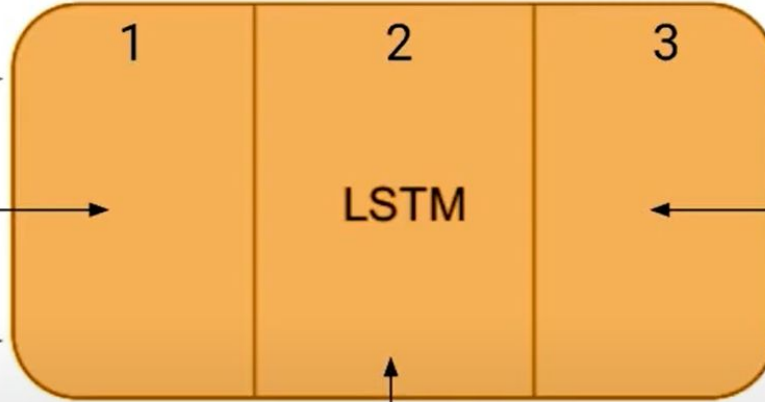
Long Term Memory

$C_{t-1}$

Forget irrelevant  
information

$H_{t-1}$

Short Term Memory



add/update new  
information

Pass updated  
information

# LSTM Architecture Experimentation

## Parameters Tested:

1. Number of epochs
2. Input sequence length
3. Number of LSTM Units
4. Dropout Rate
5. Number of LSTM Stacks

## Procedure:

1. Randomly select three stocks
2. Perform walk-forward validation with varying parameters
3. Parameters that achieve lowest MSE are chosen

# LSTM Final Architecture

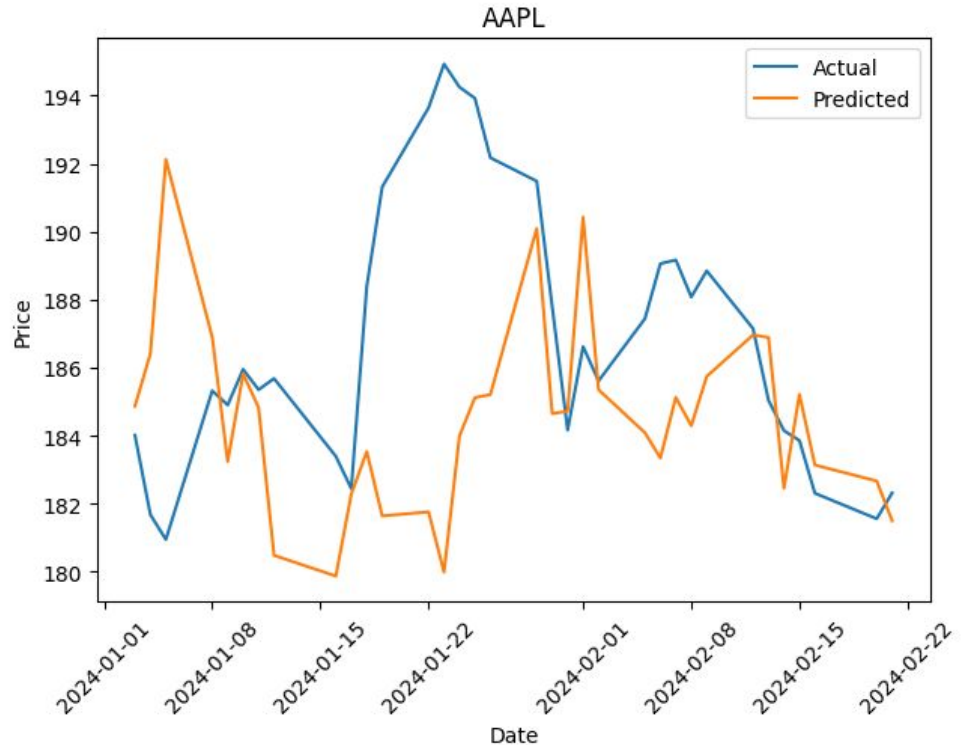
## Final Parameter Results:

1. **Epochs = 5**
2. **Input Length = 7**
3. **LSTM Units = 15 per stack**
4. **Dropout Rate = 15%**
5. **LSTM Stacks = 6**

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 7, 1)]	0
lstm (LSTM)	(None, 7, 15)	1020
dropout (Dropout)	(None, 7, 15)	0
lstm_1 (LSTM)	(None, 7, 15)	1860
dropout_1 (Dropout)	(None, 7, 15)	0
lstm_2 (LSTM)	(None, 7, 15)	1860
dropout_2 (Dropout)	(None, 7, 15)	0
lstm_3 (LSTM)	(None, 7, 15)	1860
dropout_3 (Dropout)	(None, 7, 15)	0
lstm_4 (LSTM)	(None, 7, 15)	1860
dropout_4 (Dropout)	(None, 7, 15)	0
lstm_5 (LSTM)	(None, 15)	1860
dropout_5 (Dropout)	(None, 15)	0
dense (Dense)	(None, 1)	16
Total params: 10336 (40.38 KB)		
Trainable params: 10336 (40.38 KB)		
Non-trainable params: 0 (0.00 Byte)		

# LSTM: Predicting Closing Prices

- Model trained to predict first day of 2024 (1/2/24)
- Evaluate each day in 2024:
  1. Update model with the previous day's information (1/2/24)
  2. Predict the next day's closing price (1/3/24)
- Provides more long-term view of performance

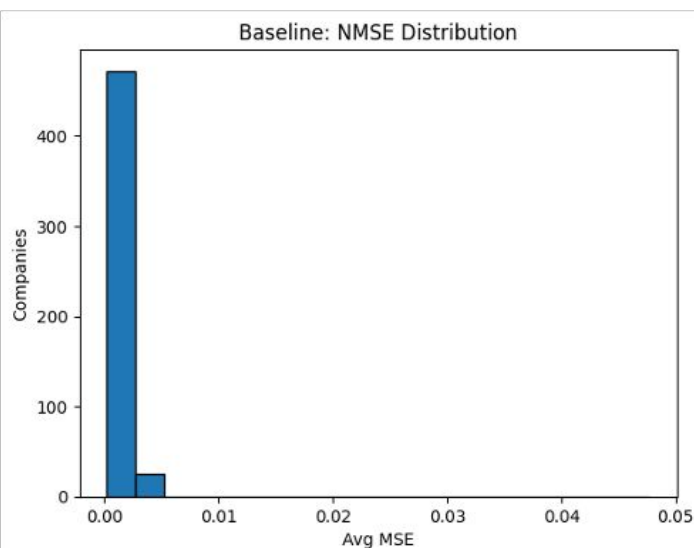
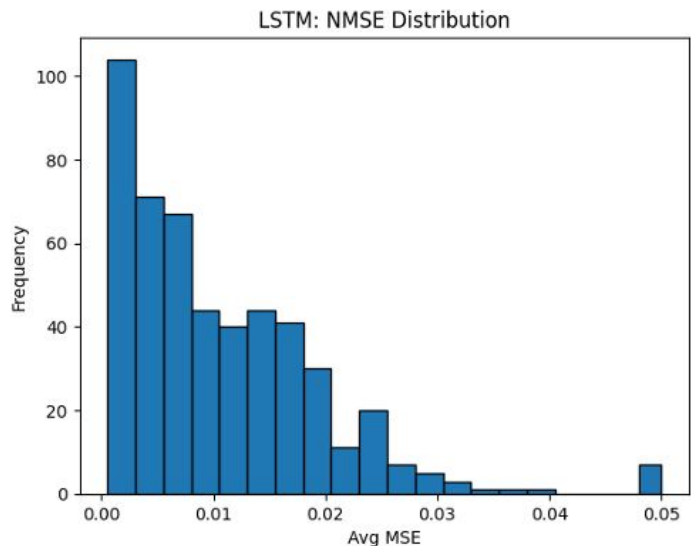


Actual vs Predicted Closing Stock Price for Apple



# LSTM: Comparing the Model to the Baseline

The baseline model of using previous day's closing price generally performs better than the LSTM models.



## Why?

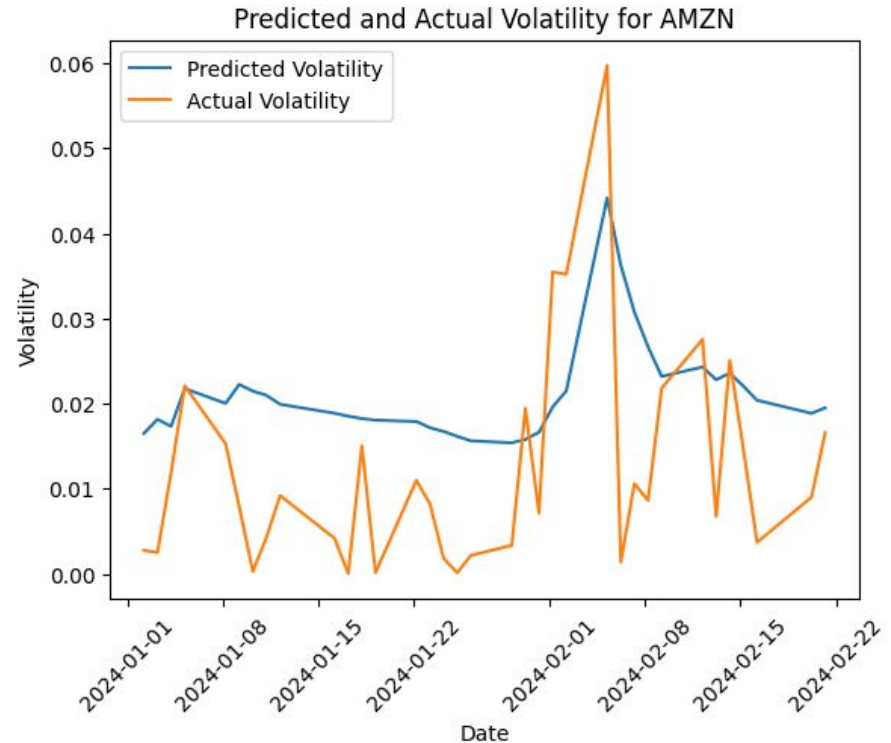
Did not use  
exogenous  
variables

Used the same  
architecture for all  
stocks

Fine-tuned  
parameters on only  
3 stocks

# GARCH: Predicting Volatility

- **Volatility** – measure of the amount of uncertainty involved in the size of changes in a stock's value
- **GARCH** – Generalized Autoregressive Conditional Heteroskedasticity
- GARCH(1,1): look at yesterday's data
- Models not fine-tuned or evaluated for scope purposes



Actual vs Predicted Closing Stock Price for Amazon

# Slot Filling and Parsing

"What 10 stocks are expected to have the highest increase in price for tomorrow?"

**Transformer Encoder\***

"o select-stock o o o o o order-by-forecasted\_price\_change-desc  
select-forecasted\_price o o o o"

**SlotParser\*\***

```
SELECT stock, forecasted_price, forecasted_price_change
FROM stock_data
ORDER BY forecasted_price_change DESC
LIMIT 10
```

\* no pre-trained embeddings or transformers were used for slot filling

\*\* SlotParser is a custom function we developed explicitly for this project

## Training Dataset

- **32** starter questions
- **25** permutations per starter from ChatGPT
- **2,365** examples in total with regex help

## Transformer Encoder

- **8,747,163** weights

## Model Evaluation

- **94%** output token accuracy
- **93%** slot-filling accuracy



# Future Work

## Forecasts: LSTM & Garch

### Use Exogenous Variables

Include data such as previous day high/low, market trends, and trading volume.

### Optimize Architecture for All Stocks Independently

Not all stocks will have the same patterns and therefore different model structures may improve performance

### Hypertune Garch Model

### Automated data pulls & model updates

## SlotMarketSQL Application

### Incorporating Filtering Capability

Allow for the application to return filtered tables based on user conditions.

### Test pre-trained embedding models (GloVe, BERT)

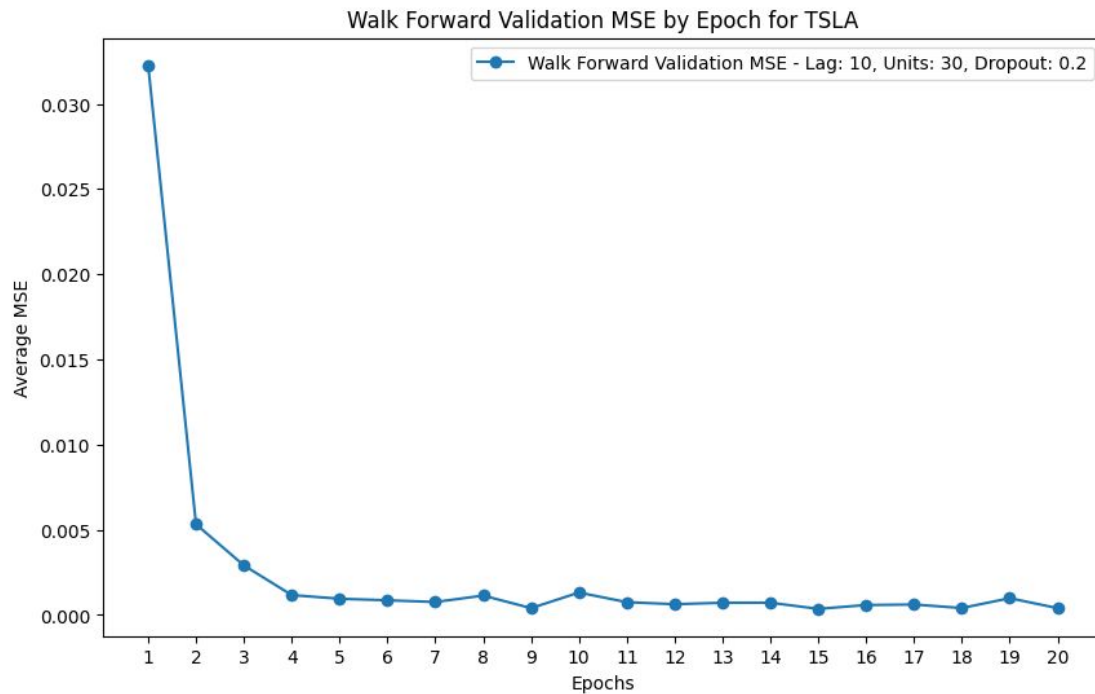
### Consider Using LLMs

Remove the dependency to hand-label queries for training & be able to handle more complex user requests.

Questions?

Backup Slides

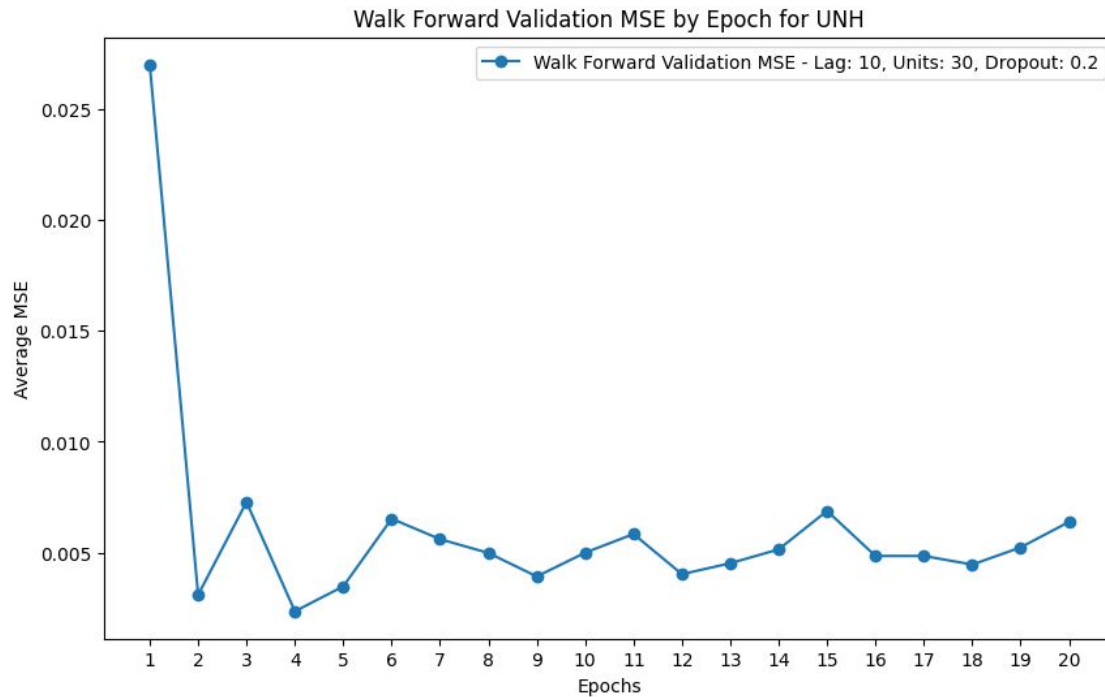
# Backup Slides



*Figure 1: Average mean squared error from walk-forward validation across varying numbers of epochs for TSLA*

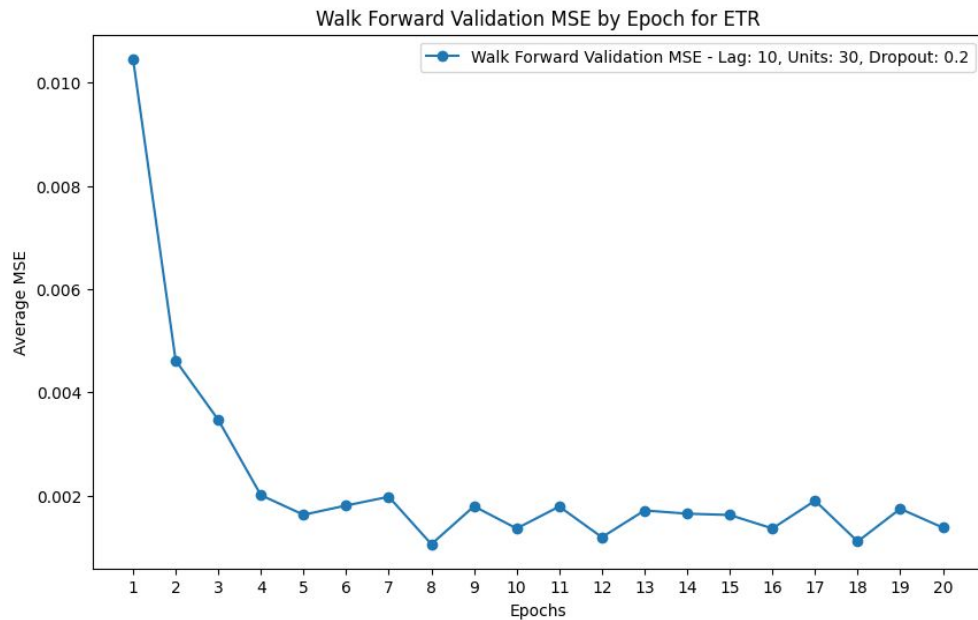


# Backup Slides



*Figure 2: Average mean squared error from walk-forward validation across varying numbers of epochs for UNH*

# Backup Slides



*Figure 3: Average mean squared error from walk-forward validation across varying numbers of epochs for ETR*

# Backup Slides

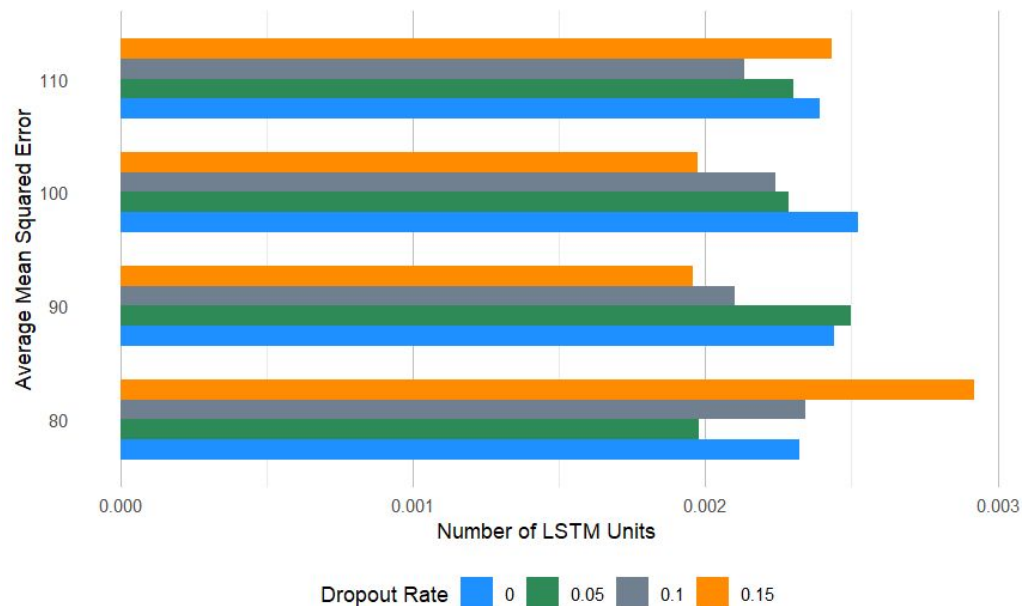
<b>Input Sequence Length</b>	<b>Average MSE [10<sup>-3</sup>]</b>
7	2.57
14	3.88
28	3.84

*Table 1: Average mean squared error from walk-forward validation across varying input sequence lengths.*

<b>Stock</b>	<b>Number of LSTM Units</b>	<b>Dropout Rate</b>
TSLA	90	10%
UNH	100	0%
ETR	50	15%

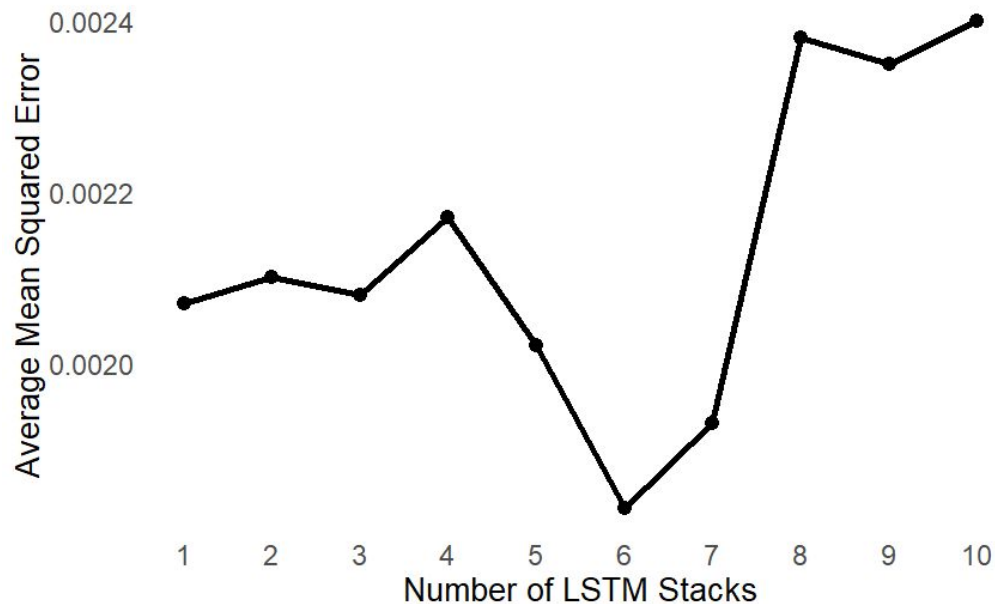
*Table 2: Number of LSTM units and dropout rate selected from Keras hyperparameter searches for each stock.*

# Backup Slides



*Figure 4: Average mean squared error from walk-forward validation across varying numbers of LSTM units and dropout rates.*

# Backup Slides



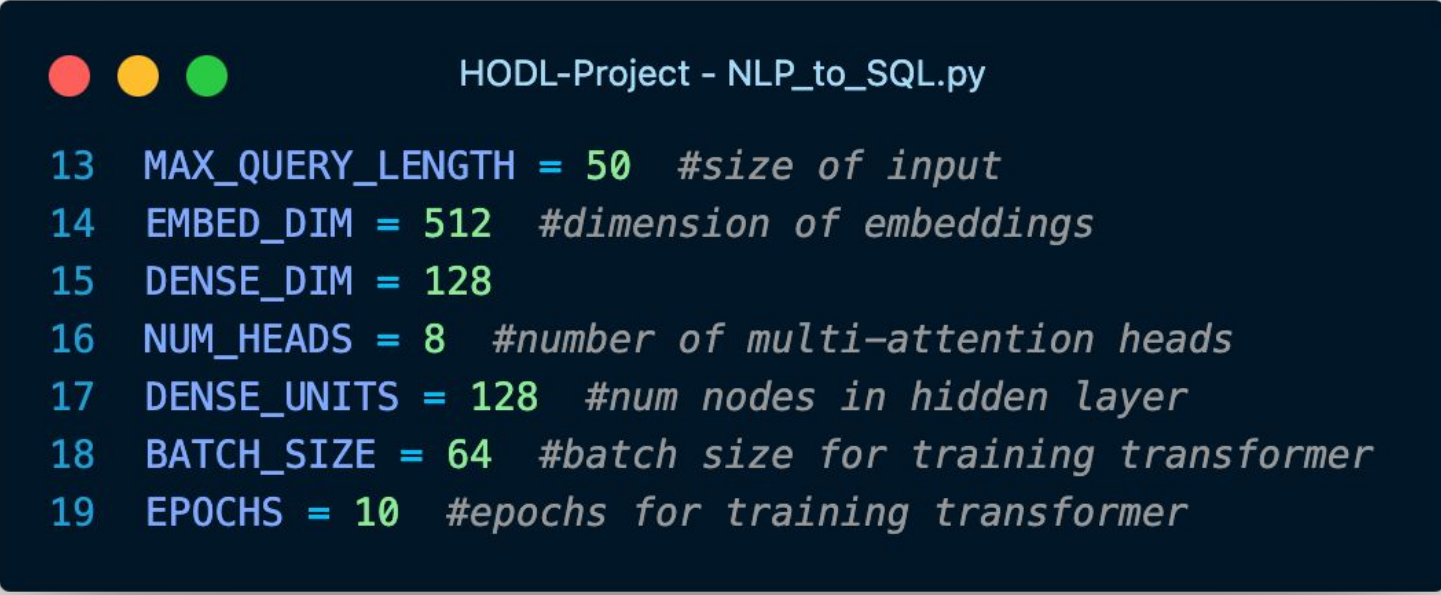
*Figure 5: Average mean squared error from walk-forward validation across varying numbers of LSTM stacks.*

# Backup Slides

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 7, 1)]	0
lstm (LSTM)	(None, 7, 15)	1020
dropout (Dropout)	(None, 7, 15)	0
lstm_1 (LSTM)	(None, 7, 15)	1860
dropout_1 (Dropout)	(None, 7, 15)	0
lstm_2 (LSTM)	(None, 7, 15)	1860
dropout_2 (Dropout)	(None, 7, 15)	0
lstm_3 (LSTM)	(None, 7, 15)	1860
dropout_3 (Dropout)	(None, 7, 15)	0
lstm_4 (LSTM)	(None, 7, 15)	1860
dropout_4 (Dropout)	(None, 7, 15)	0
lstm_5 (LSTM)	(None, 15)	1860
dropout_5 (Dropout)	(None, 15)	0
dense (Dense)	(None, 1)	16
Total params: 10336 (40.38 KB)		
Trainable params: 10336 (40.38 KB)		
Non-trainable params: 0 (0.00 Byte)		

*Figure 6: Architecture of final models*

# Backup Slides



```
13 MAX_QUERY_LENGTH = 50 #size of input
14 EMBED_DIM = 512 #dimension of embeddings
15 DENSE_DIM = 128
16 NUM_HEADS = 8 #number of multi-attention heads
17 DENSE_UNITS = 128 #num nodes in hidden layer
18 BATCH_SIZE = 64 #batch size for training transformer
19 EPOCHS = 10 #epochs for training transformer
```

Figure 7: Hyperparameters for transformer model

# Backup Slides

```
HODL-Project - NLP_to_SQL.py

31 # CREATE VECTORIZER (QUERY & SLOTS)
32 vectorize_query_text = keras.layers.TextVectorization(
33     max_tokens=None, #no maximum vocabulary
34     output_sequence_length=MAX_QUERY_LENGTH, #pad or truncate output to value
35     output_mode="int", #vector has index of vocabulary
36     standardize="lower_and_strip_punctuation", #convert input to lowercase and rmv punctuation
37     split="whitespace", #split values based on whitespace
38     ngrams=1 #only look at whole words
39 )
40 vectorize_slot_text = keras.layers.TextVectorization(
41     max_tokens=None, #no maximum vocabulary
42     output_sequence_length=MAX_QUERY_LENGTH,
43     output_mode="int", #vector has index of vocabulary
44     standardize="lower", #convert input to lowercase [can't do punctuation b/c of dashes]
45     split="whitespace", #split values based on whitespace
46     ngrams=1 #only look at whole words
47 )
48
49 # CREATE VOCABULARY AND VECTORIZED TRAINING DATA
50 vectorize_query_text.adapt(train_query) #build vocabulary
51 query_train = vectorize_query_text(train_query) #vectorized training queries
52 query_test = vectorize_query_text(test_query) #vectorized testing queries
53 QUERY_VOCAB_SIZE = vectorize_query_text.vocabulary_size() #total vocabulary of queries
54
55 vectorize_slot_text.adapt(train_slotfilling) #build slot vocabulary
56 slots_train = vectorize_slot_text(train_slotfilling) #vectorized training slots
57 slots_test = vectorize_slot_text(test_slotfilling) #vectorized testing slots
58 SLOT_VOCAB_SIZE = vectorize_slot_text.vocabulary_size() #total vocabulary of slots
```

Figure 8: Tokenization Layers



# Backup Slides

```
HODL-Project - NLP_to_SQL.py

60 # BUILD KERAS MODEL
61 inputs = keras.Input(shape=(MAX_QUERY_LENGTH,))
62 embedding = PositionalEmbedding(MAX_QUERY_LENGTH,
63                                QUERY_VOCAB_SIZE,
64                                EMBED_DIM)
65 x = embedding(inputs)
66 encoder_out = TransformerEncoder(EMBED_DIM,
67                                  DENSE_DIM,
68                                  NUM_HEADS)(x)
69 x = keras.layers.Dense(DENSE_UNITS, activation="relu", name="Dense_Layer")(encoder_out)
70 x = keras.layers.Dropout(0.25, name="Dropout_Layer")(x)
71 outputs = keras.layers.Dense(SLOT_VOCAB_SIZE, activation="softmax", name="Softmax_Layer")(x)
72
73 model = keras.Model(inputs, outputs)
74 print()
75 print(model.summary())
76 print()
77
78 # TRAIN KERAS MODEL
79 model.compile(optimizer="adam",
80               loss="sparse_categorical_crossentropy",
81               metrics=["sparse_categorical_accuracy"])
82 history = model.fit(query_train, slots_train,
83                     batch_size=BATCH_SIZE,
84                     epochs=EPOCHS)
85
86 # OUT-OF-SAMPLE TESTING
87 model.evaluate(query_test, slots_test)
88
89 # SAVE MODEL
90 filename = 'sql_transformer.keras'
91 model.save(filename)
92 ZipFile('model_save.zip', mode='w').write(filename)
```

Figure 9: Building keras model